

# DataBase programming course

## 2023

Milestone one of the project (“ Indexing”)

Made by:  
**Mohamed Bassem 2003731**  
**Mina Ehab 2005830**

[14/11/2023]

---

## Query 1

Index created on	index type
Composer	B+tree

### Index Usage

Scan = 1	Read = 978	Fetch = 0
----------	------------	-----------

### Index implementation

Create index composer\_Btree on "Track"("Composer");

### Justification

B-tree indexes are well-suited for equality queries, and in this case, finding rows where "Composer" is null is essentially an equality condition. .And index is created on table track on column Composer to speed up the retrieval of rows that match the condition

Query plan Time	Query execution Time
6996.934496999756	250157.12619600317

Latency	TPS
Before = 0.969 ms	Before = 1032.317338
After = 0.800 ms	After = 1249.397877

### Conclusion

B-tree index on the "Composer" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 2

### Index created on

### index type

Composer

B-tree

### Index Usage

Scan = 0

Read = 0

Fetch = 0

### Index implementation

```
CREATE INDEX Composer_btree ON "Track"("Composer");
```

### Justification

The Index on the Composer column did not work. I tried many other columns and it did not work. The best index for this query that it always uses is its primary key and no other column works to be an index for this query. I also tried different indexing methods but still didn't work.

### Query plan Time

### Query execution Time

6749.7116190000515

609971.7622899954

### Latency

### TPS

Before = 2.247 ms

Before = 445.038683

After = 2.202 ms

After = 440.389110

### Conclusion

The values of Latency and Tps is not changed as the index I created is not used and in the after and before results the index that was used is the primary key of the table

### Query 3

Index created on	index type
Composer	B+tree

#### Index Usage

Scan = 1	Read = 8	Fetch = 0
----------	----------	-----------

#### Index implementation

```
CREATE INDEX Composer_btree ON "Track"("Composer")
```

#### Justification

The B-tree index on the "Composer" column is effective for the query's filtering condition (WHERE t."Composer" = 'AC/DC'). The index allows for efficient lookups based on equality conditions. And index is created on table track on column Composer to speed up the retrieval of rows that match the condition

Query plan Time	Query execution Time
12291.082295000024	75586.67088399908

Latency	TPS
Before = 1.279 ms	Before = 781.711950
After = 0.987ms	After = 1013.041451

#### Conclusion

B-tree index on the "Composer" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 4

### Index created on

### index type

Composer

Hash index

### Index Usage

Scan = 1

Read = 8

Fetch = 0

### Index implementation

```
CREATE INDEX Composer_hash ON "Track" using hash("Composer")
```

### Justification

Hash index on the "Composer" column is effective for the query's filtering condition (WHERE t."Composer" = 'AC/DC'). The index allows for efficient lookups based on equality conditions..And index is created on table track on column Composer to speed up the retrieval of rows that match the condition

### Query plan Time

### Query execution Time

13347.085638999974

78512.6090140009

### Latency

### TPS

Before = 1.378

Before = 725.757526

After = 1.245

After = 803.320963

### Conclusion

Hash index on the "Composer" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 5

### Index created on

UnitPrice

### index type

Partial index

### Index Usage

Scan = 1

Read = 213

Fetch = 213

### Index implementation

```
CREATE INDEX UNitPrice_partial_index ON "Track" ("UnitPrice") WHERE  
"UnitPrice" != 0.99;
```

### Justification

The partial index on "UnitPrice" is designed to cover the specific condition in the query (WHERE "UnitPrice" != 0.99). This allows the database to efficiently locate and retrieve rows where the "UnitPrice" is not equal to 0.99. And index is created on table track on column UnitPrice to speed up the retrieval of rows that match the condition

### Query plan Time

1712.0258639999865

### Query execution Time

57190.35292300116

### Latency

Before = 0.940 ms

After = 0.653 ms

### TPS

Before = 1064.128705

After = 1532.404431

### Conclusion

The partial index on the "UnitPrice" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 6

### Index created on

UnitPrice

### index type

Partial index

### Index Usage

Scan = 0

Read = 0

Fetch = 0

### Index implementation

```
CREATE INDEX UNitPrice_partial_index ON "Track" ("UnitPrice") WHERE  
"UnitPrice" = 0.99;
```

### Justification

The Index on the UnitPrice column did not work. I tried many other columns and it did not work. The best index for this query that it always uses is its primary key and no other column works to be an index for this query. I also tried different indexing methods but still didn't work.

### Query plan Time

131842.88444800072

### Query execution Time

326208.58733399806

### Latency

Before = 2.694 ms

After = 2.630 ms

### TPS

Before = 371.263929

After = 365.467290

### Conclusion

The values of Latency and Tps is not changed as the index I created is not used and in the after and before results the index that was used is the primary key of the table

## Query 7

Index created on	index type
AlbumId	B+tree

### Index Usage

Scan = 4	Read = 4	Fetch = 0
----------	----------	-----------

### Index implementation

```
CREATE INDEX Q7_AlbumId ON "Track" ("AlbumId");
```

### Justification

The B+tree index type was chosen for the "AlbumId" index in the "Track" table. B+tree is well-suited for range queries and provides efficient traversal of data in a sorted order, which aligns with the typical access patterns for querying data by AlbumId. Having an index on table Track on column AlbumId can speed up the join operations.

Query plan Time	Query execution Time
132324.90745900074	327274.50591099774

Latency	TPS
Before =7.768 ms	Before = 128.727539
After = 7.326	After = 136.495015

### Conclusion

The marginal decrease in latency and the increase TPS suggest that the index implementation improved the efficiency of queries involving the "AlbumId" column, contributing to a more responsive system.



## Query 8

### Index created on

ArtistId

### index type

B+ tree

### Index Usage

Scan = 2

Read = 2

Fetch = 2

### Index implementation

```
-- CREATE INDEX Q8_artistid ON "Album" ("ArtistId");
```

### Justification

I went with a B+tree index for the "ArtistId" column in my "Album" table because it suits my query. It's perfect for sorting and quickly finding data within a range, which is crucial for joins and conditions on "ArtistId." This choice makes my query more efficient, speeding up data retrieval based on specific conditions.

### Query plan Time

442826.9468029906

### Query execution Time

1393653.2967069934

### Latency

Before = 0.819 ms

After = 0.461 ms

### TPS

Before = 1221.389305

After = 2170.206100

### Conclusion

In summary, the utilisation of a B+tree index on the "ArtistId" column in the "Album" table resulted in a notable decrease in query execution latency and a substantial increase in TPS. This suggests that the index implementation significantly improved the efficiency of queries related to "ArtistId," leading to a more responsive and performant system.

## Query 9

### Index created on

AlbumId

### index type

B+ tree

### Index Usage

Scan = 6

Read = 8

Fetch = 4

### Index implementation

```
CREATE INDEX Q9_AlbumId ON "Track" ("AlbumId");
```

### Justification

I decided to use a B+tree index for the "AlbumId" column in my "Track" table to boost query performance. B+tree indexes are great for range queries, which fits well with how I usually access data when identifying albums. This choice is about optimizing the way I query data based on album identification, making it more efficient.

### Query plan Time

55329.78263099972

### Query execution Time

150704.62613399967

### Latency

Before = 1.835 ms

After = 1.552 ms

### TPS

Before = 545.049249

After = 700.267436

### Conclusion

In summary, the implementation of the B+tree index on the "AlbumId" column in the "Track" table resulted in a reduced query execution latency and an increase in throughput. This suggests that the index optimization positively impacted the efficiency of queries involving the "AlbumId" column, contributing to overall improved system performance.

## Query 10

Index created on	index type
ArtistId	B+tree

### Index Usage

Scan = 2	Read = 2	Fetch = 2
----------	----------	-----------

### Index implementation

```
-- CREATE INDEX Q10_Artist ON "Album" ("ArtistId")
```

### Justification

I added a B+tree index to the "ArtistId" column in my "Album" table to improve query performance. B+tree indexes excel at handling range queries, which is exactly what I need for queries related to artist identification. Having an index on table Album on column ArtistId can speed up the join operations.

Query plan Time	Query execution Time
57166.80287899973	154677.46051399963

Latency	TPS
Before = 6.055 ms	Before = 165.154535
After = 5.573 ms	After = 179.443129

### Conclusion

B-tree index on the "ArtistId" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency.

## Query 11

### Index created on

### index type

TrackId

B+tree

### Index Usage

Scan = 52

Read = 149

Fetch = 145

### Index implementation

```
CREATE INDEX PlayListTrack_btree ON "PlaylistTrack"("TrackId")
```

### Justification

B-tree is suitable for equality conditions. Having an index on table PlaylistTrack on column TrackId can speed up the join operations.

### Query plan Time

### Query execution Time

58969.88314600196

294841.7196140023

### Latency

### TPS

Before = 1.608 ms

Before = 621.861520

After = 1.044 ms

After = 957.952080

### Conclusion

B-tree index on the "TrackId" column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 12

Index created on	index type
BillingCountry	B+tree

### Index Usage

Scan = 1	Read = 35	Fetch =0
----------	-----------	----------

### Index implementation

```
CREATE INDEX BillingCountry_btree ON "Invoice"("BillingCountry")
```

### Justification

B-tree index on the " BillingCountry " column is effective for the query's filtering condition (where "BillingCountry" = 'Brazil'). The index allows for efficient lookups based on equality conditions. nd index is created on table Invoice on column BillingCountry to speed up the retrieval of rows that match the condition

Query plan Time	Query execution Time
5142.190808000033	6977.268794000108

Latency	TPS
Before = 0.424 ms	Before = 2358.554231
After = 0.303 ms	After = 3304.933925

### Conclusion

B+tree index on the " BillingCountry " column has significantly improved the performance of the query by increasing the TPS and reducing the average latency

## Query 13

### Index created on

Country

### index type

B-tree

### Index Usage

Scan = 0

Read = 0

Fetch = 0

### Index implementation

```
CREATE INDEX Country_btree ON "Customer"("Country")
```

### Justification

The Index on the Country column did not work. I tried many other columns and it did not work. The best index for this query that it always uses is its primary key and no other column works to be an index for this query. I also tried different indexing methods but still didn't work.

### Query plan Time

1662.9098509999894

### Query execution Time

1459.2451340000346

### Latency

Before = 0.530 ms

After = 0.510 ms

### TPS

Before = 1886.852231

After = 1876.557902

### Conclusion

The values of Latency and Tps is not changed as the index I created is not used and in the after and before results the index that was used is the primary key of the table

## Query 14

### Index created on

### index type

SupportRepld

B-tree

### Index Usage

Scan = 0

Read = 0

Fetch = 0

### Index implementation

```
CREATE INDEX customerId_btree ON "Customer" ("SupportRepld");
```

### Justification

The Index on the SupportRepld column did not work. I tried many other columns and it did not work. The best index for this query that it always uses is its primary key and no other column works to be an index for this query. I also tried different indexing methods but still didn't work.

### Query plan Time

### Query execution Time

9826.950984999976

6479.058575000022

### Latency

### TPS

Before = 0.308 ms

Before = 3241.903061

After = 0.315 ms

After = 3244.807140

### Conclusion

The values of Latency and Tps is not changed as the index I created is not used and in the after and before results the index that was used is the primary key of the tables

## Query 15

### Index created on

### index type

BillingCity

B-tree

### Index Usage

Scan = 0

Read = 0

Fetch = 0

### Index implementation

Create Index BillingCity\_btree on "Invoice" ("BillingCity ")

### Justification

The Index on the BillingCity column did not work. I tried many other columns and it did not work. The best index for this query that it always uses is its primary key and no other column works to be an index for this query. I also tried different indexing methods but still didn't work.

### Query plan Time

### Query execution Time

132999.68133400058

328611.7562689977

### Latency

### TPS

Before = 0.682 ms

Before = 1466.002395

After = 0.641

After = 1456.231735

### Conclusion

The values of Latency and Tps is not changed as the index I created is not used and in the after and before results the index that was used is the primary key of the table



## Query 16

### Index created on

### index type

Total

B+tree

### Index Usage

Scan = 1

Read = 62

Fetch = 62

### Index implementation

```
-- CREATE INDEX Q16_total ON "Invoice" ("Total");
```

### Justification

I chose to use a B+tree index on the "Total" column in the "Invoice" table to enhance query performance, especially for operations involving total amounts. B+tree indexes are ideal for range queries and efficiently navigating sorted numerical data. And index is created on table Invoice on column Total to speed up the retrieval of rows that match the condition.

### Query plan Time

### Query execution Time

57693.12498899976

155905.4141099997

### Latency

### TPS

Before = 0.722 ms

Before = 1384.303326

After = 0.347 ms

After = 2885.865868

### Conclusion

In summary, the implementation of the B+tree index on the "Total" column in the "Invoice" table resulted in a substantial reduction in query execution latency and a significant increase in TPS.

## Query 17

### Index created on

### index type

("InvoiceId", "TrackId", "UnitPrice", "Quantity")

B+tree

### Index Usage

Scan = 2

Read = 2

Fetch = 0

### Index implementation

```
-- CREATE INDEX Q17_Track ON "InvoiceLine" ("InvoiceId", "TrackId") INCLUDE ("UnitPrice", "Quantity");
```

### Justification

I created a composite index on ("InvoiceId", "TrackId") in the "InvoiceLine" table, including ("UnitPrice", "Quantity") as additional columns. This choice aims to optimize query performance. Composite indexes are valuable for addressing multiple columns in queries, and including extra columns improves the index's capacity to fulfill query needs without requiring additional lookups.

### Query plan Time

### Query execution Time

58476.09782999976

157562.39142599975

### Latency

### TPS

Before = 1.474 ms

Before = 678.271825

After = 1.127 ms

After = 887.199125

### Conclusion

In summary, the creation of the composite index on ("InvoiceId", "TrackId") with included columns ("UnitPrice", "Quantity") in the "InvoiceLine" table resulted in a notable reduction in query execution latency and an increase in TPS.

## Query 18

### Index created on

### index type

Invoiceld

B+tree

### Index Usage

Scan = 9

Read = 40

Fetch = 38

### Index implementation

```
-- CREATE INDEX Q18_Invoiceld ON "InvoiceLine"("Invoiceld");
```

### Justification

I added a B+tree index to the "Invoiceld" column in the "InvoiceLine" table to boost query performance, especially when filtering or joining based on the invoice identifier.

### Query plan Time

### Query execution Time

59216.41521299976

159395.2489249999

### Latency

### TPS

Before = 0.591 ms

Before = 1693.470540

After = 0.466 ms

After = 2144.660304

### Conclusion

In summary, the creation of the B+tree index on the "Invoiceld" column in the "InvoiceLine" table resulted in a notable reduction in query execution latency and an increase in TPS.